

# Project Description

The goal of the project is to conduct the research to help optimize marketing expenses, based on the data from the analytical department at Yandex.Afisha for the period from June 2017 through May 2018.

## Available data:

- Server logs with data on Yandex.Afisha visits from June 2017 through May 2018
- Dump file with all orders for the period
- Marketing expenses statistics

## Objects of study:

- How people use the product
- When they start to buy
- How much money each customer brings
- When they pay off

## Download the data and prepare it for analysis

```
In [1]: home = %pwd
print(home)
if home != 'C:/Users/Coami':

    !pip install -Uq matplotlib --user
    !pip install -Uq numpy --user
    !pip install -Uq pandas --user
    !pip install -Uq plotly --user

    !pip install -Uq seaborn --user
    !pip install -Uq sidetable --user
```

C:\Users\Sophie\Personal\_proj

```
WARNING: The script f2py.exe is installed in 'C:\Users\Sophie\AppData\Roaming\Python\Python39\Scripts' which is not on PATH.
Consider adding this directory to PATH or, if you prefer to suppress this warning, use --no-warn-script-location.
ERROR: pip's dependency resolver does not currently take into account all the packages that are installed. This behaviour is the source of the following dependency conflicts.
daal4py 2021.5.0 requires daal==2021.4.0, which is not installed.
scipy 1.7.3 requires numpy<1.23.0,>=1.16.5, but you have numpy 1.23.4 which is incompatible.
numba 0.55.1 requires numpy<1.22,>=1.18, but you have numpy 1.23.4 which is incompatible.
```

```
In [2]: #importing libraries

import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
import datetime
import random

#!pip install -Uq seaborn --user
#!pip install -Uq plotly --user
```

```
# !pip install plotly
# !pip install -Uq plotly --user
import sys
#!conda install --yes --prefix {sys.prefix} plotly
import plotly as px
import plotly.express as px
```

## Optimize the data for analysis.

```
In [4]: try:
        visits = pd.read_csv('visits_log_us.csv', nrows=500)
    except:
        visits = pd.read_csv('/datasets/visits_log_us.csv', nrows=500)

    visits.info(memory_usage='deep')
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 500 entries, 0 to 499
Data columns (total 5 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Device      500 non-null    object
1   End Ts      500 non-null    object
2   Source Id   500 non-null    int64
3   Start Ts    500 non-null    object
4   Uid         500 non-null    uint64
dtypes: int64(1), object(3), uint64(1)
memory usage: 113.1 KB
```

```
In [5]: visits['Device'].value_counts()
```

```
Out[5]: desktop    363
        touch      137
        Name: Device, dtype: int64
```

```
In [6]: visits['Source Id'].value_counts()
```

```
Out[6]: 4      159
        3      122
        5       92
        2       52
        1       39
        9       19
        10      17
        Name: Source Id, dtype: int64
```

```
In [7]: try:
        visits = pd.read_csv(
            'visits_log_us.csv',
            dtype={'Device': 'category', 'Source Id': 'category'},
            parse_dates=['Start Ts', 'End Ts'])
    except:
        visits = pd.read_csv(
            '/datasets/visits_log_us.csv',
            dtype={'Device': 'category', 'Source Id': 'category'},
            parse_dates=['Start Ts', 'End Ts'])

    visits=visits.rename(columns={"Device": "device", "End Ts": "end_ts", "Source Id": "source_id"})
    visits.info(memory_usage='deep')
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 359400 entries, 0 to 359399
Data columns (total 5 columns):
#   Column      Non-Null Count  Dtype
---  -
```

```

-----
0    device      359400 non-null    category
1    end_ts      359400 non-null    datetime64[ns]
2    source_id   359400 non-null    category
3    start_ts    359400 non-null    datetime64[ns]
4    uid         359400 non-null    uint64
dtypes: category(2), datetime64[ns](2), uint64(1)
memory usage: 8.9 MB

```

```

In [8]: try:
        orders = pd.read_csv('/datasets/orders_log_us.csv', nrows=500)

except:
    orders = pd.read_csv('orders_log_us.csv', nrows=500)
orders.info(memory_usage='deep')

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 500 entries, 0 to 499
Data columns (total 3 columns):
#   Column      Non-Null Count  Dtype
---  -
0    Buy Ts      500 non-null    object
1    Revenue     500 non-null    float64
2    Uid         500 non-null    uint64
dtypes: float64(1), object(1), uint64(1)
memory usage: 45.0 KB

```

```

In [9]: orders.head()

```

```

Out[9]:

```

	Buy Ts	Revenue	Uid
0	2017-06-01 00:10:00	17.00	10329302124590727494
1	2017-06-01 00:25:00	0.55	11627257723692907447
2	2017-06-01 00:27:00	0.37	17903680561304213844
3	2017-06-01 00:29:00	0.55	16109239769442553005
4	2017-06-01 07:58:00	0.37	14200605875248379450

```

In [10]: try:
        orders = pd.read_csv(
            '/datasets/orders_log_us.csv',
            parse_dates=['Buy Ts'],)

except:
    orders = pd.read_csv(
        'orders_log_us.csv',
        parse_dates=['Buy Ts'],)

orders=orders.rename(columns={"Buy Ts": "buy_ts", "Revenue": "revenue", "Uid": "uid"})

orders.info(memory_usage='deep')

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50415 entries, 0 to 50414
Data columns (total 3 columns):
#   Column      Non-Null Count  Dtype
---  -
0    buy_ts      50415 non-null    datetime64[ns]
1    revenue     50415 non-null    float64
2    uid         50415 non-null    uint64
dtypes: datetime64[ns](1), float64(1), uint64(1)
memory usage: 1.2 MB

```

```
In [11]: try:
costs = pd.read_csv('costs_us.csv', nrows=500)
except:
costs = pd.read_csv('/datasets/costs_us.csv', nrows=500)

costs.info(memory_usage='deep')
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 500 entries, 0 to 499
Data columns (total 3 columns):
#   Column      Non-Null Count  Dtype
---  -
0   source_id    500 non-null    int64
1   dt           500 non-null    object
2   costs        500 non-null    float64
dtypes: float64(1), int64(1), object(1)
memory usage: 40.7 KB
```

```
In [12]: costs['source_id'].value_counts()
```

```
Out[12]: 1    363
2    137
Name: source_id, dtype: int64
```

```
In [13]: costs['costs'].describe()
```

```
Out[13]: count    500.000000
mean       65.715440
std        35.374315
min         5.800000
25%        41.390000
50%        59.345000
75%        81.542500
max       272.590000
Name: costs, dtype: float64
```

```
In [14]: try:
costs = pd.read_csv(
    '/datasets/costs_us.csv',
    dtype={'source_id': 'category'},
    parse_dates=['dt'],
)
except:
costs = pd.read_csv(
    'costs_us.csv',
    dtype={'source_id': 'category'},
    parse_dates=['dt'],
)

costs.info(memory_usage='deep')
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2542 entries, 0 to 2541
Data columns (total 3 columns):
#   Column      Non-Null Count  Dtype
---  -
0   source_id    2542 non-null    category
1   dt           2542 non-null    datetime64[ns]
2   costs        2542 non-null    float64
dtypes: category(1), datetime64[ns](1), float64(1)
memory usage: 43.0 KB
```

## Conclusion

The data was uploaded.

All dates within the data were changed to the date datatype. Columns were renamed for convenience. The source and devices information was read as categorical datatype due to the memory limitations.

## Make reports and calculate metrics

### Product

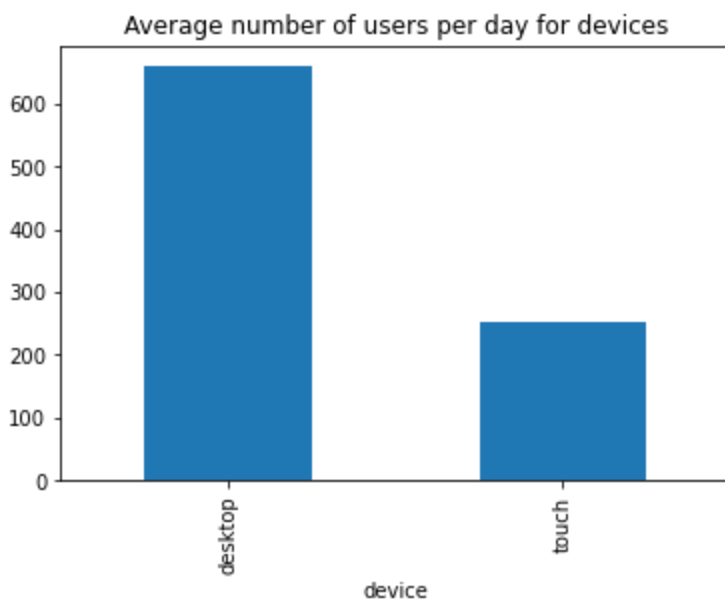
How many people use it every day, week, and month?

```
In [15]: visits['session_year'] = visits['start_ts'].dt.year
visits['session_month'] = visits['start_ts'].dt.month
visits['session_week'] = visits['start_ts'].dt.week
visits['session_date'] = visits['start_ts'].dt.date
mau_total = (
    visits.groupby(['session_year', 'session_month'])
    .agg({'uid': 'nunique'})
    .mean()
)
dau_total = (
    visits.groupby(['session_year', 'session_date'])
    .agg({'uid': 'nunique'})
    .mean()
)
wau_total = (
    visits.groupby(['session_year', 'session_week'])
    .agg({'uid': 'nunique'})
    .mean()
)
print(int(dau_total), 'people used it every day')
print(int(wau_total), 'people used it every week')
print(int(mau_total), 'people used it every month')
```

```
C:\Users\Sophie\AppData\Local\Temp\ipykernel_4048\3328251206.py:3: FutureWarning: Series.dt.weekofyear and Series.dt.week have been deprecated. Please use Series.dt.isocalendar().week instead.
```


```
visits['session_week'] = visits['start_ts'].dt.week
907 people used it every day
5716 people used it every week
23228 people used it every month
```

```
In [16]: visits.pivot_table(index='session_date', columns='device', values='uid', aggfunc='nunique')
plt.title('Average number of users per day for devices')
plt.show()
```




Number of the desktop version users are significantly exceed the number of mobile version users.

```
In [17]: dau = visits.pivot_table(index=('session_year', 'session_date'), values='uid', aggfunc=pd.S
fig, ax = plt.subplots(figsize=(17, 10))
sns.lineplot(data=dau_, x="session_date", y="uid") # hue="platform", markers="o")
plt.title('DAU')
plt.grid()
plt.show()
```




We can see the rise starting from September with noticeable peaks and outliers in the data.

```
In [18]: wau = visits.groupby([visits['start_ts'].astype('datetime64[W]')]).agg({'uid': 'nunique'
fig, ax = plt.subplots(figsize=(17, 10))
sns.lineplot(data=wau_, x="start_ts", y="uid") # hue="platform", markers="o")
plt.title('WAU')
plt.grid()
plt.show()
```



The graph correlates with the DAU graph, we see rise at autumn with significant peaks.

```
In [19]: visits['session_year_month'] = visits['start_ts'].dt.strftime('%Y-%m')
mau_ = visits.pivot_table(index='session_year_month', values='uid', aggfunc=pd.Series.nunique)
mau_
fig, ax = plt.subplots(figsize=(17, 10))
sns.lineplot(data=mau_, x="session_year_month", y="uid") # hue="platform", markers="o"
plt.grid()
plt.title('MAU')
plt.show()
```



The overall tendencies are that the most visitors are in November-December, and after the number of visitors start to decline.

```
In [20]: sticky_wau=dau_total/wau_total*100
sticky_mau=dau_total/mau_total*100
print('Sticky factor for week:',int(sticky_wau),'%')
print('Sticky factor for month:',int(sticky_mau),'%')
```

```
Sticky factor for week: 15 %
Sticky factor for month: 3 %
```

```
In [21]: dau_['session_date']=dau_['session_date'].astype('datetime64')
dau_['session_year_month'] = dau_['session_date'].dt.strftime('%Y-%m')

sticky_mau_=dau_.merge(mau_, how = 'left', left_on='session_year_month', right_on='sessi
sticky_mau_['sticky_f']=round(sticky_mau_['uid_x']/sticky_mau_['uid_y']*100)
sticky_mau_
```


```
Out[21]:
```

	session_year	session_date	uid_x	session_year_month	uid_y	sticky_f
0	2017	2017-06-01	605	2017-06	13259	5.0
1	2017	2017-06-02	608	2017-06	13259	5.0
2	2017	2017-06-03	445	2017-06	13259	3.0
3	2017	2017-06-04	476	2017-06	13259	4.0
4	2017	2017-06-05	820	2017-06	13259	6.0
...	...	...	...	...	...	...
359	2018	2018-05-27	620	2018-05	20701	3.0
360	2018	2018-05-28	1039	2018-05	20701	5.0
361	2018	2018-05-29	948	2018-05	20701	5.0
362	2018	2018-05-30	1289	2018-05	20701	6.0
363	2018	2018-05-31	1997	2018-05	20701	10.0

364 rows × 6 columns

```
In [22]: fig,ax=plt.subplots(figsize=(17,10))
sns.lineplot(data=sticky_mau_, x="session_date", y="sticky_f")
plt.title('Sticky factor metrics for MAU ratio')
plt.grid()
plt.show()
```





How many sessions are there per day?

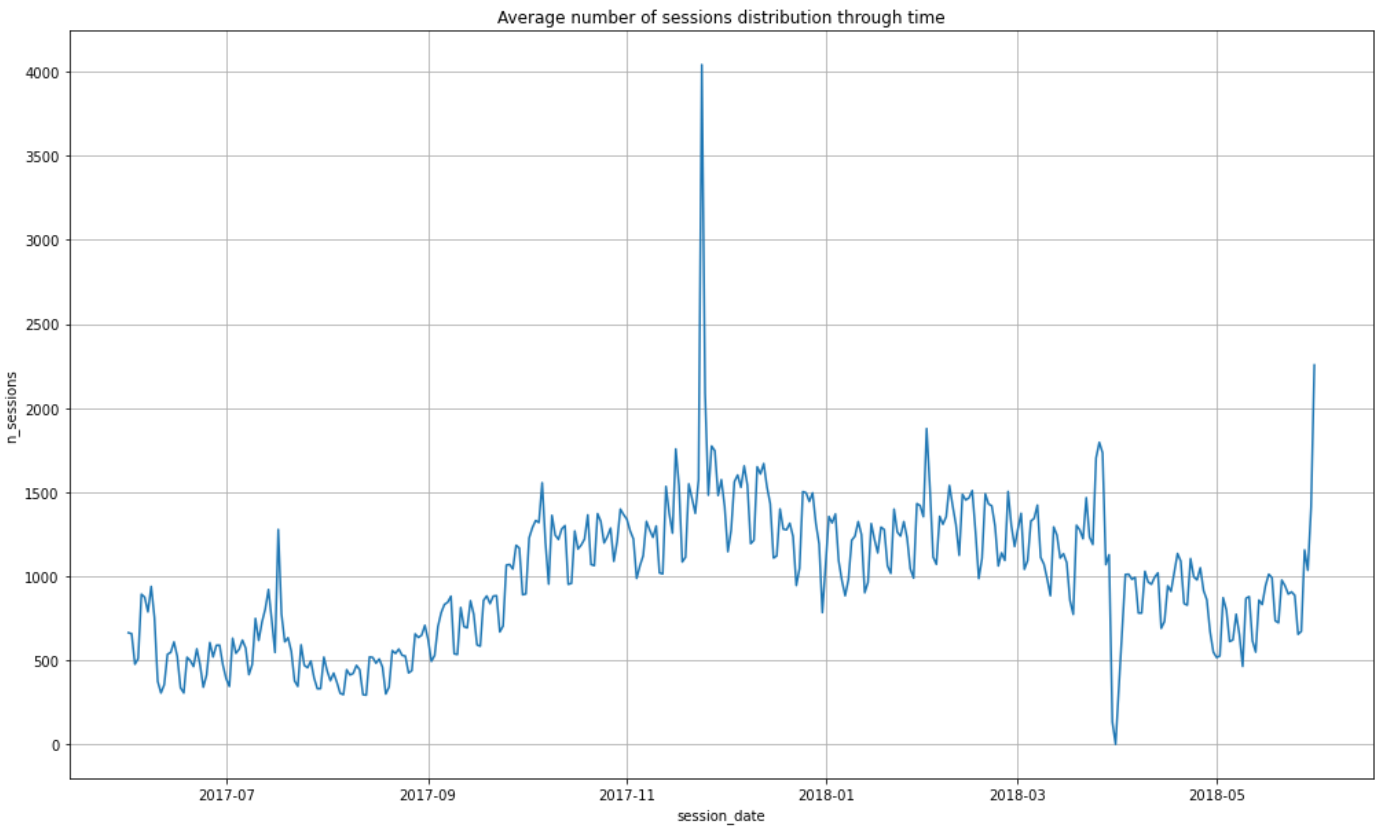
```
In [23]: sessions_per_user = visits.groupby(['session_date']).agg(
        {'uid': ['count', 'nunique']}
    ).reset_index()
    sessions_per_user.columns = ['session_date', 'n_sessions', 'n_users']
    sessions_per_user['sessions_per_user'] = (
        sessions_per_user['n_sessions'] / sessions_per_user['n_users']
    )

    print(sessions_per_user)
```

	session_date	n_sessions	n_users	sessions_per_user
0	2017-06-01	664	605	1.097521
1	2017-06-02	658	608	1.082237
2	2017-06-03	477	445	1.071910
3	2017-06-04	510	476	1.071429
4	2017-06-05	893	820	1.089024
..	...	...	...	...
359	2018-05-27	672	620	1.083871
360	2018-05-28	1156	1039	1.112608
361	2018-05-29	1035	948	1.091772
362	2018-05-30	1410	1289	1.093871
363	2018-05-31	2256	1997	1.129695

[364 rows x 4 columns]

```
In [24]: fig,ax=plt.subplots(figsize=(17,10))
    sns.lineplot(data=sessions_per_user, x="session_date", y="n_sessions")
    plt.title('Average number of sessions distribution through time')
    plt.grid()
    plt.show()
```



**Number of sessions correlates with number of users. Rise in November-December.**

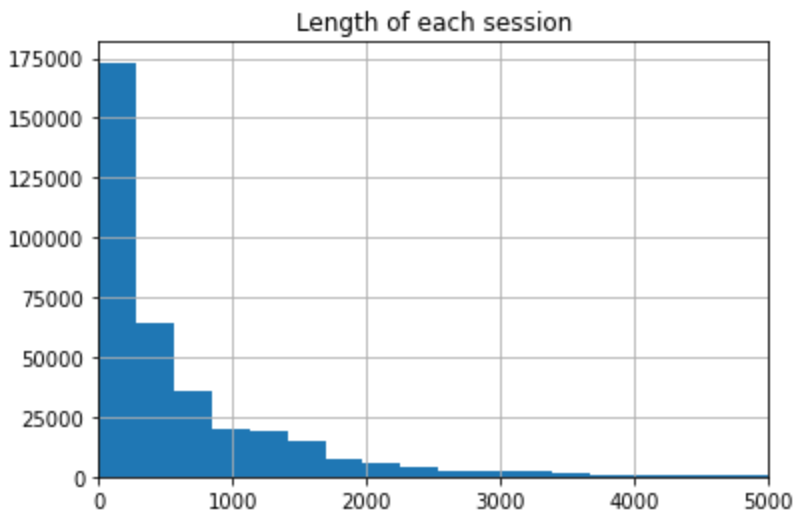
```
In [25]: print("The average number of sessions per day: {0:.2f} ".format(sessions_per_user['n_sessions_per_user']))
```

The average number of sessions per day: 987.36

**What is the length of each session?**

```
In [26]: visits['session_duration_sec'] = (
    visits['end_ts'] - visits['start_ts']
).dt.seconds
visits['session_duration_sec'].hist(bins=300)
asl=visits['session_duration_sec'].mode()[0]
print('Average session length:',asl)
plt.title('Length of each session')
plt.xlim(0,5000)
plt.show()
```

Average session length: 60



**We have unnaturally long sessions in the data, but the average session is around 1 minute.**

How often do users come back?

```
In [27]: user_activity=visits[['uid','source_id','device','start_ts']]
first_activity_date = user_activity.groupby(['uid'])['start_ts'].min()
first_activity_date.name = 'first_activity_date'
user_activity = user_activity.join(first_activity_date, on='uid')
user_activity['activity_month']= user_activity['start_ts'].dt.strftime('%Y-%m')
user_activity['first_activity_month']= user_activity['first_activity_date'].dt.strftime(
user_activity
```

Out[27]:

	uid	source_id	device	start_ts	first_activity_date	activity_month	first_activity_mo
0	16879256277535980062	4	touch	2017-12-20 17:20:00	2017-12-20 17:20:00	2017-12	2017
1	104060357244891740	2	desktop	2018-02-19 16:53:00	2018-02-19 16:53:00	2018-02	2018
2	7459035603376831527	5	touch	2017-07-01 01:54:00	2017-07-01 01:54:00	2017-07	2017
3	16174680259334210214	9	desktop	2018-05-20 10:59:00	2018-03-09 20:05:00	2018-05	2018
4	9969694820036681168	3	desktop	2017-12-27 14:06:00	2017-12-27 14:06:00	2017-12	2017
...	...	...	...	...	...	...	...
359395	18363291481961487539	2	desktop	2017-07-29 19:07:00	2017-07-29 19:07:00	2017-07	2017
359396	18370831553019119586	1	touch	2018-01-25 17:38:00	2018-01-25 17:38:00	2018-01	2018
359397	18387297585500748294	4	desktop	2018-03-03 10:12:00	2018-03-03 10:12:00	2018-03	2018
359398	18388616944624776485	5	desktop	2017-11-02 10:12:00	2017-11-02 10:12:00	2017-11	2017
359399	18396128934054549559	2	touch	2017-09-10 13:13:00	2017-09-10 13:13:00	2017-09	2017

359400 rows × 7 columns

```
In [28]: user_activity['activity_month']=pd.to_datetime(user_activity['activity_month'])
user_activity['first_activity_month']=pd.to_datetime(user_activity['first_activity_month'])
user_activity.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 359400 entries, 0 to 359399
Data columns (total 7 columns):
#   Column              Non-Null Count  Dtype
---
```

```

0    uid          359400 non-null uint64
1    source_id      359400 non-null category
2    device         359400 non-null category
3    start_ts       359400 non-null datetime64[ns]
4    first_activity_date 359400 non-null datetime64[ns]
5    activity_month  359400 non-null datetime64[ns]
6    first_activity_month 359400 non-null datetime64[ns]
dtypes: category(2), datetime64[ns](4), uint64(1)
memory usage: 14.4 MB

```

```

In [29]: user_activity['cohort_lifetime'] = (
          user_activity['activity_month'] - user_activity['first_activity_month']
        )
user_activity['cohort_lifetime'] = user_activity[
    'cohort_lifetime'
] / np.timedelta64(1, 'M')
user_activity['cohort_lifetime'] = user_activity['cohort_lifetime'].round().astype(int)

user_activity

```

```

Out[29]:

```

	uid	source_id	device	start_ts	first_activity_date	activity_month	first_activity_mo
0	16879256277535980062	4	touch	2017-12-20 17:20:00	2017-12-20 17:20:00	2017-12-01	2017-12
1	104060357244891740	2	desktop	2018-02-19 16:53:00	2018-02-19 16:53:00	2018-02-01	2018-02
2	7459035603376831527	5	touch	2017-07-01 01:54:00	2017-07-01 01:54:00	2017-07-01	2017-07
3	16174680259334210214	9	desktop	2018-05-20 10:59:00	2018-03-09 20:05:00	2018-05-01	2018-03
4	9969694820036681168	3	desktop	2017-12-27 14:06:00	2017-12-27 14:06:00	2017-12-01	2017-12
...	...	...	...	...	...	...	...
359395	18363291481961487539	2	desktop	2017-07-29 19:07:00	2017-07-29 19:07:00	2017-07-01	2017-07
359396	18370831553019119586	1	touch	2018-01-25 17:38:00	2018-01-25 17:38:00	2018-01-01	2018-01
359397	18387297585500748294	4	desktop	2018-03-03 10:12:00	2018-03-03 10:12:00	2018-03-01	2018-03
359398	18388616944624776485	5	desktop	2017-11-02 10:12:00	2017-11-02 10:12:00	2017-11-01	2017-11
359399	18396128934054549559	2	touch	2017-09-10 13:13:00	2017-09-10 13:13:00	2017-09-01	2017-09

359400 rows × 8 columns

```

In [30]: cohorts = user_activity.groupby(['first_activity_month', 'cohort_lifetime']).agg({'uid':

```

```

initial_users_count =cohorts[cohorts['cohort_lifetime']== 0][
    ['first_activity_month', 'uid']
]

initial_users_count =initial_users_count.rename(
    columns={'uid': 'cohort_users'})


cohorts =cohorts.merge(initial_users_count, on='first_activity_month')

cohorts['retention'] = cohorts['uid'] / cohorts['cohort_users']

retention_pivot = cohorts.pivot_table(
    index='first_activity_month',
    columns='cohort_lifetime',
    values='retention',
    aggfunc='sum',
)
retention_pivot
sns.set(style='white')
plt.figure(figsize=(13, 9))
plt.title('Cohorts: User Retention')
sns.heatmap(
    retention_pivot, annot=True, fmt='.1%', linewidths=1, linecolor='black', vmin= 0, vm
)

```

Out[30]: <AxesSubplot:title={'center':'Cohorts: User Retention'}, xlabel='cohort\_lifetime', ylabel='first\_activity\_month'>



## Conclusion

1. Based on the average amount of visitors throughout the year we can see that visits start to increase in autumn with the peak on November-December, and then slowly drops till reaches the minimum in summer with occasional peaks. We can assume that it's connected with the seasons

- of the events, for example, most theatre productions open in autumn, but summer usually is the "dead" season. Also, there are usually a lot of holidays during autumn and winter.
- 2. Each user has one session per day on average. The dynamic of how many sessions per day is coincidental with the DAU.
- 3. We have anomaly long sessions within the data. The median session is 1 minute.
- 4. The highest retention belongs to the June cohort. A lot of people use the platform and stick to it through the autumn and winter. As it was mentioned before, it might be because events are seasonal.

## Sales

### When do people start buying?

```
In [31]: orders_=orders.pivot_table(index='uid',values='buy_ts', aggfunc = 'min').reset_index()
orders_.head()
orders_.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 36523 entries, 0 to 36522
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  -
0    uid      36523 non-null   uint64
1   buy_ts   36523 non-null   datetime64[ns]
dtypes: datetime64[ns](1), uint64(1)
memory usage: 570.8 KB
```

```
In [32]: visits_=visits.pivot_table(index='uid',values='start_ts', aggfunc = 'min').reset_index()
visits_.head(10)
visits_.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 228169 entries, 0 to 228168
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  -
0    uid      228169 non-null   uint64
1   start_ts  228169 non-null   datetime64[ns]
dtypes: datetime64[ns](1), uint64(1)
memory usage: 3.5 MB
```

```
In [33]: buys=orders_.merge(visits_[['uid','start_ts']], on='uid')
buys['conversion']=buys['buy_ts'] - buys['start_ts']
buys['conversion'] = buys[
    'conversion'
] / np.timedelta64(1, 'D')
buys['conversion'] = buys['conversion'].round().astype(int)

buys
```

```
Out[33]:
```

	uid	buy_ts	start_ts	conversion
0	313578113262317	2018-01-03 21:51:00	2017-09-18 22:49:00	107
1	1575281904278712	2017-06-03 10:13:00	2017-06-03 10:13:00	0
2	2429014661409475	2017-10-11 18:33:00	2017-10-11 17:14:00	0
3	2464366381792757	2018-01-28 15:54:00	2018-01-27 20:10:00	1
4	2551852515556206	2017-11-24 10:14:00	2017-11-24 10:14:00	0


...	...	...	...	...
36518	18445147675727495770	2017-11-24 09:03:00	2017-08-20 13:30:00	96
36519	18445407535914413204	2017-09-22 23:55:00	2017-09-22 23:48:00	0
36520	18445601152732270159	2018-03-26 22:54:00	2017-08-07 11:51:00	231
36521	18446156210226471712	2018-02-18 19:34:00	2017-11-07 10:01:00	103
36522	18446167067214817906	2017-10-17 10:16:00	2017-10-17 10:05:00	0

36523 rows × 4 columns

```
In [34]: import plotly.express as px
```

```
In [35]: fig = px.histogram(buys, x="conversion", nbins=30, title='Conversion time')
fig.show()
```

## Conversion time



Most of the orders were made within the short time after the usage of the website.

```
In [36]: buys['conversion'].median()
```

```
Out[36]: 0.0
```

```
In [37]: print(' The overall conversion is {:.1%}'.format(buys['uid'].nunique()/visits_['uid'].nu
The overall conversion is 16.0%
```

## How many orders do they make during a given period of time?

```
In [38]: #defining cohort month through first purchase
orders['month']=orders['buy_ts'].astype('datetime64[M]')
first_order_date = orders.groupby(['uid'])['buy_ts'].min()
first_order_date.name = 'first_order_date'
orders=orders.join(first_order_date, on='uid')
orders['first_order_month'] = orders['first_order_date'].astype('datetime64[M]')

orders.head()
```

```
Out[38]:
```

	buy_ts	revenue	uid	month	first_order_date	first_order_month
0	2017-06-01 00:10:00	17.00	10329302124590727494	2017-06-01	2017-06-01 00:10:00	2017-06-01
1	2017-06-01 00:25:00	0.55	11627257723692907447	2017-06-01	2017-06-01 00:25:00	2017-06-01
2	2017-06-01 00:27:00	0.37	17903680561304213844	2017-06-01	2017-06-01 00:27:00	2017-06-01
3	2017-06-01 00:29:00	0.55	16109239769442553005	2017-06-01	2017-06-01 00:29:00	2017-06-01
4	2017-06-01 07:58:00	0.37	14200605875248379450	2017-06-01	2017-06-01 07:58:00	2017-06-01

```
In [39]: orders.sample(5)
```

```
Out[39]:
```

	buy_ts	revenue	uid	month	first_order_date	first_order_month
33340	2018-02-07 09:54:00	3.67	1631356693099033031	2018-02-01	2018-02-07 09:54:00	2018-02-01
12254	2017-10-11 20:16:00	6.11	4148716401875468484	2017-10-01	2017-10-11 20:16:00	2017-10-01
43657	2018-04-10 12:10:00	4.89	3971030471691518392	2018-04-01	2018-04-10 12:10:00	2018-04-01
45602	2018-04-26 14:19:00	1.98	6093855166159309373	2018-04-01	2017-07-11 23:19:00	2017-07-01
28026	2018-01-04 13:44:00	2.44	13719347206112370408	2018-01-01	2018-01-04 13:44:00	2018-01-01

```
In [40]: #let's define cohort size
cohort_sizes = orders.groupby('first_order_month').agg({'uid': 'nunique'}).reset_index()
cohort_sizes.columns=['first_order_month','cohort_size']
cohort_sizes.head()
```

```
Out[40]:
```

	first_order_month	cohort_size
0	2017-06-01	2023
1	2017-07-01	1923
2	2017-08-01	1370
3	2017-09-01	2581
4	2017-10-01	4340



```

In [41]: #calculating number of purchases for cohort and month
cohort=orders.groupby(['first_order_month', 'month'])['revenue'].count().reset_index()
cohort.columns=['first_order_month', 'month', 'orders']
#merge cohort with cohort size
cohort=cohort.merge(cohort_sizes, on=['first_order_month'])
cohort['age_month'] = ((cohort['month'] - cohort['first_order_month']) / np.timedelta64(
cohort['orders_per_buyer']=cohort['orders']/cohort['cohort_size']
cohort.head()

```

```

Out[41]:
   first_order_month  month  orders  cohort_size  age_month  orders_per_buyer
0      2017-06-01  2017-06-01    2354         2023         0.0         1.163618
1      2017-06-01  2017-07-01     177         2023         1.0         0.087494
2      2017-06-01  2017-08-01     174         2023         2.0         0.086011
3      2017-06-01  2017-09-01     226         2023         3.0         0.111715
4      2017-06-01  2017-10-01     292         2023         4.0         0.144340

```

```

In [42]: cohort_piv=cohort.pivot_table(
        index='first_order_month',
        columns='age_month',
        values='orders_per_buyer',
        aggfunc='sum'
    ).cumsum(axis=1)

cohort_piv.round(2).fillna('')

```

```

Out[42]:
   age_month  0.0  1.0  2.0  3.0  4.0  5.0  6.0  7.0  8.0  9.0  10.0  11.0
first_order_month
2017-06-01  1.16  1.25  1.34  1.45  1.59  1.7  1.84  1.92  2.03  2.1  2.15  2.19
2017-07-01  1.14  1.19  1.25  1.31  1.34  1.39  1.42  1.44  1.47  1.49  1.51
2017-08-01  1.12  1.2  1.27  1.33  1.39  1.44  1.47  1.53  1.56  1.6
2017-09-01  1.14  1.22  1.28  1.35  1.37  1.42  1.46  1.48  1.5
2017-10-01  1.14  1.22  1.25  1.28  1.31  1.34  1.35  1.38
2017-11-01  1.18  1.28  1.32  1.37  1.41  1.42  1.45
2017-12-01  1.15  1.21  1.26  1.3  1.32  1.34
2018-01-01  1.12  1.19  1.24  1.25  1.28
2018-02-01  1.12  1.18  1.21  1.22
2018-03-01  1.17  1.22  1.27
2018-04-01  1.10  1.18
2018-05-01  1.09
2018-06-01  1.00

```


What is the average purchase size?

```

In [43]: avg_check=orders.groupby(['uid'])['revenue'].mean().reset_index()
fig = px.histogram(avg_check, x="revenue", title='Average purchase size')
fig.show()

```

## Average purchase size



Most of the purchase sizes are within 10.

```
In [44]: avg_cohort=orders.groupby(['first_order_month', 'month'])['revenue'].mean().reset_index()
avg_cohort['age_month'] = ((avg_cohort['month'] - avg_cohort['first_order_month']) / np.
avg_cohort.head()
```

```
Out[44]:
```

	first_order_month	month	revenue	age_month
0	2017-06-01	2017-06-01	4.060106	0.0
1	2017-06-01	2017-07-01	5.547006	1.0
2	2017-06-01	2017-08-01	5.088161	2.0
3	2017-06-01	2017-09-01	8.545575	3.0
4	2017-06-01	2017-10-01	7.084178	4.0

```
In [45]: avg_cohort_piv=avg_cohort.pivot_table(
    index='first_order_month',
    columns='age_month',
    values='revenue',
    aggfunc='mean'
)

avg_cohort_piv.round(2).fillna('')
```

```
Out[45]:
```

	age_month	0.0	1.0	2.0	3.0	4.0	5.0	6.0	7.0	8.0	9.0	10.0	11.0
first_order_month													
2017-06-01		4.06	5.55	5.09	8.55	7.08	6.83	6.97	6.76	5.28	8.01	12.04	6.04

2017-07-01	5.29	6.45	9.99	6.64	4.72	3.66	3.79	5.45	5.35	11.79	5.65
2017-08-01	4.72	5.99	6.28	6.62	7.96	6.27	5.89	7.11	8.7	5.6	
2017-09-01	4.97	13.17	8.35	62.57	15.43	15.32	16.77	11.21	7.79		
2017-10-01	4.37	7.41	5.13	5.59	5.1	5.07	4.28	4.01			
2017-11-01	4.37	4.1	4.47	6.28	4.44	3.73	4.6				
2017-12-01	4.11	4.23	20.07	26.08	15.95	14.11					
2018-01-01	3.69	4.44	6.45	7.52	2.71						
2018-02-01	3.71	4.58	3.45	3.87							
2018-03-01	4.14	5.97	6.33								
2018-04-01	4.25	6.2									
2018-05-01	4.29										
2018-06-01	3.42										

```
In [46]: orders['revenue'].mean()
```

```
Out[46]: 4.999646930476993
```

How much money do they bring? (LTV)

```
In [47]: orders.head(10)
```

```
Out[47]:
```

	buy_ts	revenue	uid	month	first_order_date	first_order_month
0	2017-06-01 00:10:00	17.00	10329302124590727494	2017-06-01	2017-06-01 00:10:00	2017-06-01
1	2017-06-01 00:25:00	0.55	11627257723692907447	2017-06-01	2017-06-01 00:25:00	2017-06-01
2	2017-06-01 00:27:00	0.37	17903680561304213844	2017-06-01	2017-06-01 00:27:00	2017-06-01
3	2017-06-01 00:29:00	0.55	16109239769442553005	2017-06-01	2017-06-01 00:29:00	2017-06-01
4	2017-06-01 07:58:00	0.37	14200605875248379450	2017-06-01	2017-06-01 07:58:00	2017-06-01
5	2017-06-01 08:43:00	0.18	10402394430196413321	2017-06-01	2017-06-01 08:43:00	2017-06-01
6	2017-06-01 08:54:00	1.83	12464626743129688638	2017-06-01	2017-06-01 08:54:00	2017-06-01
7	2017-06-01 09:22:00	1.22	3644482766749211722	2017-06-01	2017-06-01 09:22:00	2017-06-01
8	2017-06-01 09:22:00	3.30	17542070709969841479	2017-06-01	2017-06-01 09:22:00	2017-06-01
9	2017-06-01 09:23:00	0.37	1074355127080856382	2017-06-01	2017-06-01 09:23:00	2017-06-01

```
In [48]: #get the revenue per cohort in each month
ltv_cohort=orders.groupby(['first_order_month','month'])['revenue'].sum().reset_index()
ltv_cohort.columns = ['first_order_month','month','revenue']
```

```
#merge with the cohort size
ltv_cohort=ltv_cohort.merge(cohort_sizes,on=['first_order_month'])
ltv_cohort['age']=((ltv_cohort['month'] - ltv_cohort['first_order_month']) / np.timedelta64(1,'D')).astype(int)
ltv_cohort['ltv']=ltv_cohort['revenue']/ltv_cohort['cohort_size']
ltv_cohort
```


Out[48]:

	first_order_month	month	revenue	cohort_size	age	ltv
0	2017-06-01	2017-06-01	9557.49	2023	0.0	4.724414
1	2017-06-01	2017-07-01	981.82	2023	1.0	0.485329
2	2017-06-01	2017-08-01	885.34	2023	2.0	0.437637
3	2017-06-01	2017-09-01	1931.30	2023	3.0	0.954671
4	2017-06-01	2017-10-01	2068.58	2023	4.0	1.022531
...	...	...	...	...	...	...
74	2018-03-01	2018-05-01	1114.87	3533	2.0	0.315559
75	2018-04-01	2018-04-01	10600.69	2276	0.0	4.657597
76	2018-04-01	2018-05-01	1209.92	2276	1.0	0.531599
77	2018-05-01	2018-05-01	13925.76	2988	0.0	4.660562
78	2018-06-01	2018-06-01	3.42	1	0.0	3.420000

79 rows × 6 columns

In [49]:

```
ltv_cohort_piv=ltv_cohort.pivot_table(
    index='first_order_month',
    columns='age',
    values='ltv',
    aggfunc='sum'
).cumsum(axis=1)
plt.figure(figsize=(13, 9))
ltv_cohort_piv.index=ltv_cohort_piv.index.astype(str)
sns.heatmap(ltv_cohort_piv, annot=True, fmt='.2f', linewidths=1, linecolor='grey', cbar_
).set(title='LTV')
plt.figure(figsize=(13, 9))
plt.show()
```



<Figure size 936x648 with 0 Axes>


## Conclusion

1. People start buying the same day they start using the service. The overall conversion is 16%.
2. On average people make one or two purchases each month.
3. The average purchase size is around 5. The most profitable time is at the end of the year, it correlates with all that was mentioned above, autumn months and December are the most profitable. Also, people seem to spend most money several months after using the service.
4. The most profitable cohort are September one followed by summer users.

## Marketing


How much money was spent? Overall/per source/over time

```
In [50]: costs_=costs.sort_values(by=['dt','source_id'])
costs_.head()
costs_dt=costs_.groupby(['dt']).sum().reset_index()
fig,ax=plt.subplots(figsize=(17,10))
sns.lineplot(data=costs_dt, x="dt", y="costs")
plt.title('Marketing costs over time')
plt.grid()
plt.show()
```



The most money on the marketing was spent in November-December.

```
In [51]: costs_s=costs_.groupby(['source_id']).sum().reset_index()
fig,ax=plt.subplots(figsize=(17,10))
sns.lineplot(data=costs_s, x="source_id", y="costs")
plt.title('Marketing costs per source')
plt.grid()
plt.show()
```



Sources 3,4 and 5 are the most expensive.

```
In [52]: print('Total marketing cost is {}'.format(costs['costs'].sum()))
```

Total marketing cost is 329131.62

How much did customer acquisition from each of the sources cost?

```
In [53]: orders
```

Out[53]:

	buy_ts	revenue	uid	month	first_order_date	first_order_month
0	2017-06-01 00:10:00	17.00	10329302124590727494	2017-06-01	2017-06-01 00:10:00	2017-06-01
1	2017-06-01 00:25:00	0.55	11627257723692907447	2017-06-01	2017-06-01 00:25:00	2017-06-01
2	2017-06-01 00:27:00	0.37	17903680561304213844	2017-06-01	2017-06-01 00:27:00	2017-06-01
3	2017-06-01 00:29:00	0.55	16109239769442553005	2017-06-01	2017-06-01 00:29:00	2017-06-01
4	2017-06-01 07:58:00	0.37	14200605875248379450	2017-06-01	2017-06-01 07:58:00	2017-06-01
...	...	...	...	...	...	...
50410	2018-05-31 23:50:00	4.64	12296626599487328624	2018-05-01	2018-05-31 23:50:00	2018-05-01
50411	2018-05-31 23:50:00	5.80	11369640365507475976	2018-05-01	2018-05-31 23:50:00	2018-05-01
50412	2018-05-31 23:54:00	0.30	1786462140797698849	2018-05-01	2018-05-31 23:54:00	2018-05-01
50413	2018-05-31 23:56:00	3.67	3993697860786194247	2018-05-01	2018-05-31 23:56:00	2018-05-01
50414	2018-06-01 00:02:00	3.42	83872787173869366	2018-06-01	2018-06-01 00:02:00	2018-06-01

50415 rows × 6 columns

```
In [54]: costs['costs_month']=costs['dt'].astype('datetime64[M]')
costs
```

Out[54]:

	source_id	dt	costs	costs_month
0	1	2017-06-01	75.20	2017-06-01
1	1	2017-06-02	62.25	2017-06-01
2	1	2017-06-03	36.53	2017-06-01
3	1	2017-06-04	55.00	2017-06-01
4	1	2017-06-05	57.08	2017-06-01
...	...	...	...	...
2537	10	2018-05-27	9.92	2018-05-01
2538	10	2018-05-28	21.26	2018-05-01
2539	10	2018-05-29	11.32	2018-05-01
2540	10	2018-05-30	33.15	2018-05-01

2542 rows × 4 columns

```
In [55]: costs_by_month=costs.groupby(['costs_month'])['costs'].sum().reset_index()
costs_by_month.head()
```

```
Out[55]:
```

	costs_month	costs
0	2017-06-01	18015.00
1	2017-07-01	18240.59
2	2017-08-01	14790.54
3	2017-09-01	24368.91
4	2017-10-01	36322.88

```
In [56]: customers_per_moth=orders.groupby(['first_order_month'])['uid'].nunique().reset_index()
customers_per_moth.columns=['costs_month','customers']
customers_per_moth.head()
```

```
Out[56]:
```

	costs_month	customers
0	2017-06-01	2023
1	2017-07-01	1923
2	2017-08-01	1370
3	2017-09-01	2581
4	2017-10-01	4340


```
In [57]: CAC_per_month=costs_by_month.merge(customers_per_moth,how='left',on=['costs_month'])
CAC_per_month['CAC']=CAC_per_month['costs']/CAC_per_month['customers']
CAC_per_month.head()
```

```
Out[57]:
```


	costs_month	costs	customers	CAC
0	2017-06-01	18015.00	2023	8.905091
1	2017-07-01	18240.59	1923	9.485486
2	2017-08-01	14790.54	1370	10.796015
3	2017-09-01	24368.91	2581	9.441654
4	2017-10-01	36322.88	4340	8.369327

```
In [58]: fig = px.line(CAC_per_month, x="costs_month", y="CAC", title='CAC')
fig.show()
```

CAC







The highest CAC was during July-August period.

### CAC per source

In [59]: visits

	device	end_ts	source_id	start_ts	uid	session_year	session_month	session_week
0	touch	2017-12-20 17:38:00	4	2017-12-20 17:20:00	16879256277535980062	2017	12	1
1	desktop	2018-02-19 17:21:00	2	2018-02-19 16:53:00	104060357244891740	2018	2	1
2	touch	2017-07-01 01:54:00	5	2017-07-01 01:54:00	7459035603376831527	2017	7	1
3	desktop	2018-05-20 11:23:00	9	2018-05-20 10:59:00	16174680259334210214	2018	5	1
4	desktop	2017-12-27 14:06:00	3	2017-12-27 14:06:00	9969694820036681168	2017	12	1
...	...	...	...	...	...	...	...	...
359395	desktop	2017-07-29 19:07:19	2	2017-07-29 19:07:00	18363291481961487539	2017	7	1
359396	touch	2018-01-25 17:38:19	1	2018-01-25 17:38:00	18370831553019119586	2018	1	1
359397	desktop	2018-03-03 10:12:19	4	2018-03-03 10:12:00	18387297585500748294	2018	3	1
359398	desktop	2017-	5	2017-	18388616944624776485	2017	11	1

		11-02		11-02				
		10:12:19		10:12:00				
		2017-		2017-				
359399	touch	09-10	2	09-10	18396128934054549559	2017	9	:
		13:13:19		13:13:00				

359400 rows × 11 columns

```
In [60]: visits.groupby(['uid'])['source_id'].nunique().head()
```

```
Out[60]: uid
11863502262781    1
49537067089222    1
297729379853735    1
313578113262317    1
325320750514679    1
Name: source_id, dtype: int64
```

```
In [61]: first_source=visits.sort_values('start_ts').groupby('uid').first()['source_id'].reset_index()
first_source.columns=['uid', 'first_source']
first_source.head()
```

```
Out[61]:
```

	uid	first_source
0	11863502262781	3
1	49537067089222	2
2	297729379853735	3
3	313578113262317	2
4	325320750514679	5

```
In [62]: purchase=orders.merge(first_source,on=['uid'],how='left')
purchase.head()
```

```
Out[62]:
```

	buy_ts	revenue	uid	month	first_order_date	first_order_month	first_source
0	2017-06-01 00:10:00	17.00	10329302124590727494	2017-06-01	2017-06-01 00:10:00	2017-06-01	1
1	2017-06-01 00:25:00	0.55	11627257723692907447	2017-06-01	2017-06-01 00:25:00	2017-06-01	2
2	2017-06-01 00:27:00	0.37	17903680561304213844	2017-06-01	2017-06-01 00:27:00	2017-06-01	2
3	2017-06-01 00:29:00	0.55	16109239769442553005	2017-06-01	2017-06-01 00:29:00	2017-06-01	2
4	2017-06-01 07:58:00	0.37	14200605875248379450	2017-06-01	2017-06-01 07:58:00	2017-06-01	3

```
In [63]: costs_by_month_source=costs.groupby(['costs_month','source_id'])['costs'].sum().reset_index()
costs_by_month_source.head()
```

```
Out[63]:
```

	costs_month	source_id	costs
0	2017-06-01	1	1125.61
1	2017-06-01	10	314.22
2	2017-06-01	2	2427.38

3	2017-06-01	3	7731.65
4	2017-06-01	4	3514.80

```
In [64]: customers_per_moth_source=purchase.groupby(['first_order_month','first_source'])['uid'].
customers_per_moth_source.columns=['costs_month','source_id','customers']
customers_per_moth_source.head()
```

```
Out[64]:
```

	costs_month	source_id	customers
0	2017-06-01	1	190
1	2017-06-01	10	95
2	2017-06-01	2	235
3	2017-06-01	3	638
4	2017-06-01	4	413


```
In [65]: CAC_per_month_source=costs_by_month_source.merge(customers_per_moth_source,how='left',on
CAC_per_month_source['CAC']=CAC_per_month_source['costs']/CAC_per_month_source['customer
CAC_per_month_source.head()
```

```
Out[65]:
```

	costs_month	source_id	costs	customers	CAC
0	2017-06-01	1	1125.61	190	5.924263
1	2017-06-01	10	314.22	95	3.307579
2	2017-06-01	2	2427.38	235	10.329277
3	2017-06-01	3	7731.65	638	12.118574
4	2017-06-01	4	3514.80	413	8.510412

```
In [66]: #plotting cac dynamics
fig = px.line(CAC_per_month_source, x="costs_month", y="CAC",color='source_id',title='CA
fig.show()
```

CAC





Sources 2 and 3 are the most expensive for the customer's acquirement. CAC for sources 9 and 10 are beneath the average.

How worthwhile where the investments? (ROI)

ROI per cohort


```
In [67]: CAC_per_month_ROI=CAC_per_month[['costs_month','CAC']]
CAC_per_month_ROI.columns=['first_order_month','CAC']
ROI=ltv_cohort.merge(CAC_per_month_ROI,on=['first_order_month'],how='left')
ROI.head()
```

```
Out[67]:
```

	first_order_month	month	revenue	cohort_size	age	ltv	CAC
0	2017-06-01	2017-06-01	9557.49	2023	0.0	4.724414	8.905091
1	2017-06-01	2017-07-01	981.82	2023	1.0	0.485329	8.905091
2	2017-06-01	2017-08-01	885.34	2023	2.0	0.437637	8.905091
3	2017-06-01	2017-09-01	1931.30	2023	3.0	0.954671	8.905091
4	2017-06-01	2017-10-01	2068.58	2023	4.0	1.022531	8.905091

```
In [68]: ROI['ROI']=ROI['ltv']/ROI['CAC']
roi_piv = ROI.pivot_table(
    index='first_order_month', columns='age', values='ROI', aggfunc='mean'
).cumsum(axis=1).round(2)
```

```
In [69]: roi_piv.index=roi_piv.index.astype(str)
plt.figure(figsize=(13, 9))
sns.heatmap(roi_piv, annot=True, fmt='.2f', linewidths=1, linecolor='grey', cbar_kws= {
    }).set(title = 'ROI per cohort')
plt.show()
```



## ROI per source

```
In [70]: ltv_per_source=purchase.groupby(['first_source'])['uid','revenue'].agg({'uid':'nunique',
ltv_per_source.columns=['source_id','customers','revenue']
ltv_per_source['ltv']=ltv_per_source['revenue']/ltv_per_source['customers']
ltv_per_source
```

C:\Users\Sophie\AppData\Local\Temp\ipykernel\_4048\1304129202.py:1: FutureWarning:  
Indexing with multiple keys (implicitly converted to a tuple of keys) will be deprecated, use a list instead.

```
Out[70]:
```

	source_id	customers	revenue	ltv
0	1	2899	31090.55	10.724577
1	10	1329	4450.33	3.348631
2	2	3506	46923.61	13.383802
3	3	10473	54511.24	5.204931
4	4	10296	56696.83	5.506685
5	5	6931	52624.02	7.592558
6	6	0	0.00	NaN
7	7	1	1.22	1.220000
8	9	1088	5759.40	5.293566

```
In [71]: roi_per_source=costs_s.merge(ltv_per_source,on=['source_id'])
roi_per_source['cac']=roi_per_source['costs']/roi_per_source['customers']
roi_per_source['romi']=roi_per_source['ltv']/roi_per_source['cac']
roi_per_source
```


source_id	costs	customers	revenue	ltv	cac	romi
-----------	-------	-----------	---------	-----	-----	------

Out[71]:

0	1	20833.27	2899	31090.55	10.724577	7.186364	1.492351
1	10	5822.49	1329	4450.33	3.348631	4.381106	0.764335
2	2	42806.04	3506	46923.61	13.383802	12.209367	1.096191
3	3	141321.63	10473	54511.24	5.204931	13.493901	0.385725
4	4	61073.60	10296	56696.83	5.506685	5.931779	0.928336
5	5	51757.10	6931	52624.02	7.592558	7.467479	1.016750
6	9	5517.49	1088	5759.40	5.293566	5.071222	1.043844

```
In [72]: fig = px.bar(roi_per_source, x='source_id', y='romi', title='ROMI')
fig.update_xaxes(type='category')
fig.show()
```

## ROMI



We are intrested on ROMI above 1: sources 1,2,4,5,9. Sources 10 and 3 are signigicantly below 1.

## Conclusion

1. Marketing costs correlate with the seasonal buying peaks. The most money is spent on the source 3. 2.The costs of marketing per customer are highest for sources 3 and 2. 3.Only expenses on sources 1,2,5 and 9 are profitable.

In [ ]:

# Conclusion

As a result of the conducted research, I would recommend focusing marketing expenses on source 1, since it's the cheapest and the most profitable one.

We also see that people prefer using a desktop version of the service. So it might be worth checking if a mobile version of the website is easy to work with and/or functioning correctly.

Overall we see that interest in our service is seasonal and strongly correlated with the number of events throughout the year. Therefore we might correct our expenses on marketing depending on the season.